

Chapter - 3

Linear Models

Aviral Janveja

1 Introduction

We have already introduced a linear model in chapter 1, the Perceptron. Let us begin here by revisiting a slightly modified and enhanced version of Perceptron that works for non-linearly separable data as well.

2 Linear Classification

Let us consider a real world data-set of hand-written digits collected from India-Post offices. It is always better to test your models on real data, to get a better understanding of how your system would actually perform in reality. Here is a sample from the data-set :




Figure 1: Real Data Example

As shown above, we have a bunch of hand-written digits, collected from postal stamps. We would like to design a model, that can decipher the digits from the given images. People can sometimes write digits in weird ways, making it difficult to understand even for a human operator. Indeed, the error rate for human operators is found to be around 2.5% and we would like to see if our machine learning model can at least equal that or maybe do better.

2.1 Input Representation

Let us begin by looking at the given input data more closely. We are given a set of grayscale images containing hand-written digits, as shown below :




Figure 2: Hand-Written Digit Example

Now, each of these images is 16×16 pixels. Meaning, each digit is represented by 256 real-number attributes. The raw-input \mathbf{x} , would therefore look like :

$$\mathbf{x} = (x_1, x_2, x_3 \dots x_{256})$$

That is a very long input to represent such a simple object. If we take this raw-input and try the Perceptron directly on it, we get too many parameters :

$$\mathbf{w} = (w_1, w_2, w_3 \dots w_{256})$$

The idea of input representation is to simplify the algorithm's life. We know that it is not about the individual pixel values, when trying to recognize a digit. We can instead extract some relevant features from the raw-input and then give those to the learning model and let it figure out the pattern.

2.2 Feature Engineering

Features are basically useful information that can be extracted from the given raw-input. For example, average pixel intensity, symmetry-score and curve-score. The digit 1 for instance, will score higher on the symmetry measure compared to a 5, whereas 5 will score higher on the intensity-score, similarly 8 will get a higher curve-score than 1 and so on. Using just these three features, significantly simplifies our input representation :

$$\mathbf{x} = (x_1, x_2, x_3)$$

Admittedly, we have lost information in this process. But chances are, most of that is irrelevant information anyway. From a generalization point of view as well, going from 256 to 3 parameters is a pretty good situation. Plotting a scatter diagram for digits 1 and 5 alongside just two features - symmetry and intensity. We get the following illustration :




Figure 3: Illustration of Features

The blue points represent 1 and the red points represent 5. The horizontal-axis represents intensity and the vertical-axis represents symmetry. Indeed the red-fives are tilted more towards the right on the horizontal-axis, corresponding to their higher intensity-score. Meanwhile the blue-ones are higher on the vertical-axis owing to their higher symmetry-score. Already, just by using these two features, we see that the above data is almost linearly separable.

2.3 Pocket Perceptron

3 Linear Regression

4 References

1. CalTech Machine Learning Course - CS156, Lecture 2.