

Chapter - 3

Linear Models

Aviral Janveja

1 Introduction

We have already covered one linear model in chapter 1 - The Perceptron. Let us begin here by revisiting a slightly modified and enhanced version of Perceptron that works for non-linearly separable data as well.

2 Linear Classification

2.1 Input Representation

Feature Engineering.

2.2 Pocket Perceptron

Works for non-linearly separable data.

3 Linear Regression

4 References

1. CalTech Machine Learning Course - CS156, Lecture 2.